THE COMPLETE SIMPLE-ENGLISH STUDY GUIDE

Everything About Your Research Explained in Simple, Clear Words Without Confusing Jargon

Presenter: Beknan Chemeda | *Project: Amharic Byte Mamba & Hayyuu AI*

1. The Big Story: What is this project about in 1 minute?

💡 **THE SIMPLE SUMMARY YOU CAN SAY TO ANYONE:** Imagine you are trying to teach a computer to speak Amharic. Today's big AI tools (like ChatGPT) were built for English. When they try to read Amharic, they get confused: they cut single Amharic words into 7 or 8 tiny broken pieces. This makes normal models slow, heavy, and wasteful. In our research, we solved this completely! We threw away the word-cutter program and let the AI read raw computer bytes directly using a new, fast model called 'Mamba'. Our small Mamba model beat a model 5 times its size! Then, we gave it a brain memory like a human: it can learn new facts in 0.01 seconds and remembers them permanently by 'sleeping' and replaying them.

2. The 4 Problems We Solved (With Real-Life Analogies)

• **Problem 1: The 'Small Library' Problem (Low-Resource Data):** In English, AI companies download trillions of words from the internet. But for Amharic, the entire world has less than 500 million clean words online. You cannot just download more data because it doesn't exist. If you use machine-translated English text, it has bad grammar and ruins the AI. Therefore, we had to make the AI model itself much smarter and more efficient with the data we have.

• **Problem 2: The 'Broken Word' Problem (The Tokenizer / Token Tax):** Before an AI reads words, it uses a helper program called a 'tokenizer' (a word-cutter). Tokenizers were built for English words like 'apple' -> 1 token. But when it sees an Amharic word like 'ኢትዮጵያ', it gets confused because Ge'ez letters take 3 raw computer bytes each. It cuts that one word into 7 or 8 pieces! This means an Amharic speaker has to use 3 to 7 times more computer power and money than an English speaker for the exact same message. We call this unfair cost the 'Token Tax'.

• **Problem 3: The 'Dictionary Waste' Problem (Parameter Tax):** To use a normal tokenizer, the AI must build a massive dictionary table (e.g., 32,000 words). In a small 13-million parameter model, 8.2 million of those parameters (62.8% of the entire model!) were trapped just holding that dictionary list! Only 37% was left for actual thinking and grammar. That is like buying a 100-page book where 63 pages are just an empty index list.

• **Problem 4: The 'Amnesia' Problem (Catastrophic Forgetting):** Normal AI models have a huge flaw: if you teach them today's new news via standard training, their brain overwrites what they learned before. They learn the new fact, but suddenly forget basic grammar (they start repeating words and babbling). This is called Catastrophic Forgetting.

3. The 3 Solutions We Built

✅ **Solution 1: Byte-Level AI (Throwing Away the Word Cutter):** Instead of guessing how to cut Amharic words with a broken tokenizer, we threw away the tokenizer completely! Computers store everything as numbers from 0 to 255 (raw bytes). We let the AI read raw bytes directly. Now, the dictionary table has only 256 entries instead of 32,000! This means 97.7% of the model's brain is free for pure thinking and grammar.

✅ **Solution 2: Mamba SSM (A Fast Reader That Never Slows Down):** When you give raw bytes to a normal Transformer (like GPT), it becomes extremely slow because Transformer attention scales quadratically $O(L^2)$—if text is 3x longer, it takes 9x to 12x more memory. Mamba is a new architecture called a State Space Model. It scans text in a smooth straight line (Linear Time $O(L)$). It never slows down and uses 3x less GPU memory!

✅ **Solution 3: Human Brain Memory (Fast Notebook + Slow Brain + Sleep Replay):** How do humans learn news without forgetting how to speak? The human brain has two parts: (1) A fast Hippocampus (like a quick notepad for today's events) and (2) A slow Neocortex (for permanent language grammar). We copied this! Our model writes new facts to a fast memory matrix in 0.01 seconds without touching its main brain. Then, when it 'sleeps', it replays the facts 20 times faster into its permanent neural weights!

4. Plain-English Dictionary: Simple Definitions of Every Term

📖 **Bits-Per-Byte (BPB):** How we measure model performance. Think of it like measuring fuel economy in 'Liters per 100km'. You cannot compare two models using 'Perplexity' if one uses words and the other uses bytes. BPB measures the exact computer bits needed to guess one byte of text. LOWER BPB MEANS A SMARTER, BETTER MODEL.

📖 **Transformer:** The traditional AI architecture used in GPT-4. It compares every word to every other word (Self-Attention). It is great for English, but gets very slow and memory-hungry on long byte text.

📖 **Mamba (Selective State Space Model):** A new, modern AI architecture. Instead of comparing all words to all words, it reads text in a smooth sequence and keeps a rolling mental summary. It is fast, lightweight, and uses linear computer power.

📖 **Token Tax:** The unfair computational penalty paid by non-English languages because standard tokenizers break their words into multiple pieces.

📖 **Parameter Allocation Tax:** The percentage of an AI model's brain capacity wasted on holding a dictionary list instead of thinking layers.

📖 **Hebbian Learning:** "Neurons that fire together, wire together." A simple biological formula that lets the computer remember a new fact in 0.01 seconds by connecting numbers directly, without running slow training calculations.

📖 **Synaptic Consolidation (Sleep Replay):** What the brain does at night: taking the short-term facts from the day and replaying them quickly into the permanent brain circuits so you never forget them.

📖 **SFT (Supervised Fine-Tuning):** Teaching the AI how to behave like a helpful assistant by training on question-and-answer pairs.

📖 **Prompt-Loss Masking:** Telling the AI: 'Don't grade yourself on the user's question; only grade yourself on giving the correct answer.'

📖 **Hayyuu:** The name of our living AI assistant deployed on Telegram. Named after the Oromo word for a wise elder / counselor / scholar.

5. The Results: How to Explain the Numbers Simply

Table 1: The Core Model Comparison (Lower BPB is Better)

Model Name	Input Format	Size (Params)	Dictionary Waste	Score (Bits/Byte)
Transformer (32k)	32,000 pieces	13.0 Million	62.8% (Huge waste)	1.579 BPB (Worst)
Transformer (16k)	16,000 pieces	8.9 Million	45.7% (High waste)	1.492 BPB
TinyMamba (Ours)	Raw Bytes	2.7 Million	2.3% (Minimal)	1.322 BPB (Beat 13M!)
Mamba-Base (Ours)	Raw Bytes	26.2 Million	0.4% (Pure thinking)	1.061 BPB 🏆 (Best)

How to say this to the audience: *"Look at Table 1. The Transformer model with 13 million parameters wasted 63% of its size on a dictionary and scored 1.579 BPB. Our TinyMamba model had only 2.7 million parameters (almost 5 times smaller!), wasted almost nothing on dictionaries, and scored 1.322 BPB. This proves that smaller, smarter byte models beat larger, wasteful tokenized models!"*

6. Questions You Might Get & How to Answer Without Hesitation

❓ **Question: Why didn't you just make a better Amharic tokenizer?**
👉 Simple Answer: Because Amharic grammar modifies words with prefixes and suffixes in complex ways. Even with a 32,000-word tokenizer, more than 60% of our model's weights were wasted just storing dictionary words instead of learning grammar. By removing the tokenizer, 98% of the model is dedicated to pure thinking.

❓ **Question: If bytes make the text 3 times longer, doesn't that make the AI super slow?**
👉 Simple Answer: It makes Transformers slow because Transformers compare every letter to every other letter ($O(L^2)$). But Mamba reads in a straight line ($O(L)$). Our small Mamba model processed 384 bytes at over 4 steps per second and used 3 times less GPU memory than the Transformer.

❓ **Question: How does the model learn new news on Telegram without forgetting past grammar?**
👉 Simple Answer: We copied the human brain! When someone teaches Hayyuu a new fact, it writes it to a fast memory notepad in 0.01 seconds without changing its main neural weights. When it sleeps, it replays those notes and slowly integrates them so grammar is never ruined.

❓ **Question: What is the main message of your research?**
👉 Simple Answer: You do not need billions of dollars or giant supercomputers to build high-quality AI for African languages. By using raw bytes and efficient Mamba models, anyone can train a powerful Amharic AI on a single standard graphics card (RTX 3090).