

MASTER PRESENTATION STUDY GUIDE & DEFENSE CHEAT SHEET

Byte-Level Mamba vs. Transformer for Amharic with Brain-Inspired Continual Learning (Hayyuu)

Presenter & Author: Beknan Chemeda | Project: Data- & Compute-Efficient Generative AI

1. The 30-Second Elevator Pitch (Start Your Talk With This)

💡 **30-SECOND OPENING PITCH:** "Current AI models like GPT-4 rely on subword tokenizers that charge a heavy 'Token Tax' on Ethiopian languages—splitting a single Amharic word into 7+ byte fragments and wasting up to 63% of model parameters on vocabulary embedding tables. In this research, we eliminated the tokenizer entirely by feeding raw UTF-8 bytes directly into Mamba (Selective State Space Model), which processes sequences in linear time $O(L)$ instead of quadratic $O(L^2)$. We proved that a 2.7M byte-level Mamba outperforms a 13M Transformer while consuming 3x less memory. Furthermore, to solve catastrophic forgetting, we designed a biologically inspired dual-memory system (Hayyuu) that learns new facts in 0.01 seconds and consolidates them into permanent neural weights during sleep!"

2. The Three Core Research Problems We Solved

- **Problem 1: The Amharic Data Ceiling (<500M Words):** Unlike English which has trillions of web tokens, clean digital Amharic is scarce (<500M words in the entire world). Machine-translated data introduces corrupt grammar. Therefore, AI for Amharic CANNOT scale by brute-force data collection; it requires extreme architectural efficiency.
- **Problem 2: The Token Tax & Parameter Allocation Tax:** Standard tokenizers (SentencePiece/BPE) are biased towards Latin alphabets. Ge'ez characters (Unicode U+1200 to U+137F) require 3 bytes per character in UTF-8. In standard tokenized models, 62.8% of all trainable parameters are wasted on the vocabulary embedding matrix (Embedding Tax) rather than reasoning layers.
- **Problem 3: Quadratic Attention Bottleneck $O(L^2)$:** Because Amharic byte sequences are 3.5x longer than English token sequences, standard Transformer Self-Attention becomes computationally prohibitive (memory and compute scale with sequence length squared: $O(L^2)$). Mamba scales linearly $O(L)$.
- **Problem 4: Catastrophic Forgetting in Continual Learning:** When standard LLMs are updated on new news or live facts via backpropagation, their neural weights overwrite previously learned grammatical knowledge (Stability-Plasticity Dilemma). Our brain-inspired dual-memory solves this with zero forgetting.

3. Technical Glossary: Definitions of Every Key Term

🔧 **Bits-per-Byte (BPB):** The mathematically rigorous metric used to compare language models across different tokenizers. Perplexity cannot be compared between a tokenizer with 32k pieces and a byte model (256 pieces). BPB measures the exact information bits required to compress one raw byte: $BPB = \text{CrossEntropyLoss} / \ln(2) * (\text{Tokens/Bytes})$. Lower is better.

🔧 **Mamba (Selective State Space Model - S6):** A modern neural network architecture that replaces Transformer self-attention with a continuous state-space recurrence. It processes long sequences in $O(L)$ linear time while dynamically selecting what information to remember or forget based on the current input.

🔧 **Token Tax & Fertility Ratio:** The ratio of token count to byte count or word count. In Amharic, standard tokenizers generate ~3.5 tokens per word (Fertility = 3.5), meaning an Amharic speaker pays 3.5x more computational cost and latency than an English speaker for the same meaning.

🔧 **Parameter Allocation Tax:** The percentage of a model's total parameter count dedicated exclusively to the embedding layer ($\text{vocab_size} * d_{\text{model}}$). For 32k tokenized Transformer, 62.8% of weights were trapped in

embeddings. For Byte-Level Mamba (vocab=256), only 2.3% is used for embeddings, freeing 97.7% of parameters for deep reasoning!

🔥 **Catastrophic Forgetting:** The tendency of artificial neural networks to completely and abruptly erase previously learned knowledge upon learning new information. It is the primary obstacle to lifelong continual learning in AI.

🔥 **Hebbian Learning & Anti-Hopfield Engrams:** "Neurons that fire together wire together." A biological 1-shot associative memory rule where memory weights (M) are updated via outer products of normalized key-value vectors: $\Delta M = \eta (v - M k) k^T$ without performing expensive backpropagation.

🔥 **Complementary Learning Systems (CLS):** The neuroscience theory (McClelland et al.) stating human memory relies on two collaborating structures: the fast Hippocampus (episodic, 1-shot fact learning) and the slow Neocortex (semantic, gradual generalization). We replicated this with Hebbian Fast Memory + Mamba Neocortex.

🔥 **Synaptic Consolidation (Sharp-Wave Ripple Replay):** The biological process during NREM sleep where the hippocampus replays daytime memories at high speed (20x) into neocortical circuits, integrating new facts permanently without disrupting core linguistic grammar.

🔥 **Prompt-Loss Masking (SFT):** In Supervised Fine-Tuning, setting target loss labels to -100 for the user prompt so the model is only penalized on generating the correct assistant response, preventing prompt memorization.

🔥 **Replay Buffer Anchoring (RLHF/RLAIF):** A technique that injects foundational anchor pairs (Identity, Geography, Grammar) into online reinforcement learning gradient updates to prevent 'Mode Collapse' (where a model over-optimizes a single rewarded string).

4. Experimental Results & Benchmark Comparisons

Table 1: Pre-Training Architecture Benchmark on 1.465 GB Amharic Corpus

Model Architecture	Vocab / Input	Total Params	Embed Tax (%)	Val Loss (BPB)	Linear Complexity
Transformer (32k)	SentencePiece 32k	13.04 M	62.8% (Wasted)	1.5790 BPB	Quadratic $O(L^2)$ ❌
Transformer (16k)	SentencePiece 16k	8.96 M	45.7%	1.4920 BPB	Quadratic $O(L^2)$ ❌
TinyMamba (Byte)	Raw UTF-8 Bytes	2.73 M	2.3% (Efficient)	1.3220 BPB	Linear $O(L)$ ⚡
Mamba-Base (26M)	Raw UTF-8 Bytes	26.24 M	0.49%	1.0613 BPB 🏆	Linear $O(L)$ ⚡

Table 2: Continual Learning & Catastrophic Forgetting Benchmark

Learning Method	Learning Speed	1-Shot Fact Recall	Grammar Retention	Forgetting Status
Standard Backpropagation	~4.2 seconds (Slow)	40.0% (Poor)	42.1% (Degraded)	Severe Catastrophic Forgetting ❌
Hebbian Engram (Ours)	0.01 seconds (Instant)	100.0% (Perfect)	78.4% (Preserved)	Zero Interference ✅
Engram + Sleep Replay (Ours)	20x Offline Consolidation	100.0% (Permanent)	100.0% (Perfect)	Biologically Consolidated 🏆

5. The Living Agent: Hayyuu Autonomous System

To test our findings in real-world environments, we built and deployed 'Hayyuu'—a pure neural conversational AI on Telegram. Hayyuu features a tri-state living cycle: (1) Active Day Phase: interacting with Ethiopian users and harvesting native news streams; (2) 1-Shot Fact Acquisition: forming Hebbian memory engrams instantly upon receiving verified facts; and (3) Circadian Sleep Consolidation: entering NREM sleep every 30 rounds to replay facts into Mamba neocortical weights via Sharp-Wave Ripples at 20x speed.

6. Anticipated Presentation Questions & Defense Answers (Q&A Cheat Sheet)

Q1: Why did you choose Byte-Level modeling instead of simply training a larger SentencePiece tokenizer?

👉 Answer: In low-resource morphologically rich languages like Amharic, subword tokenizers create severe fragmentation and out-of-vocabulary anomalies. More importantly, large tokenizers create a massive 'Parameter Allocation Tax'—for an 8M parameter model, a 32k vocabulary consumes 62.8% of all parameters in the embedding matrix alone! Byte modeling uses a fixed 256-byte vocabulary (taking <1% of parameters), allowing 99% of model weights to be spent on neural reasoning layers.

Q2: Since byte sequences are ~3.5x longer, isn't byte-level AI too slow?

👉 Answer: That is true for Transformers due to quadratic $O(L^2)$ self-attention. However, Mamba uses Selective State Spaces (S6) with linear $O(L)$ complexity and chunked associative parallel scans. Our benchmarks prove that TinyMamba processes 384-byte sequences at 4.3 steps/second with only 4.2GB VRAM, outperforming Transformers while using 3x less GPU memory.

Q3: How do you mathematically prove that TinyMamba beat the Transformer when they use different tokenizers?

👉 Answer: We evaluated all models using Bits-per-Byte (BPB) normalized by UTF-8 byte length ($BPB = Loss / \ln(2) * N_tokens / N_bytes$). Perplexity is incomparable across tokenizers, but BPB is a universally invariant information-theoretic metric. TinyMamba achieved 1.322 BPB, strictly compressing text better than the 13M Transformer (1.579 BPB).

Q4: How does your Hebbian Engram memory prevent Catastrophic Forgetting during live deployment?

👉 Answer: Traditional fine-tuning uses gradient backpropagation, which alters global network weights and causes catastrophic forgetting. We implement a Complementary Learning System (CLS): new facts are written in 0.01s to a dedicated Hebbian fast weight matrix (M) using anti-Hopfield outer product updates without touching the slow Mamba backbone. During offline sleep cycles, these facts are replayed with background replay anchors, permanently transferring memory without linguistic drift.

Q5: What is the impact of this research for Ethiopia and African AI?

👉 Answer: This work proves that state-of-the-art AI for low-resource African languages does not require multi-billion-parameter clusters or billions of dollars. By eliminating tokenizers and adopting linear State Space Models with bio-inspired memory, highly capable generative AI can be trained on a single commodity GPU (Nvidia RTX 3090) within a few hours on native datasets.